

cjw NETWORK

Developers united in eZ Publish



Learnings from Real eZ Publish 5 Projects

Donat Fritschy - Webmanufaktur
Ekkehard Dörre - Coolscreen

cjw NETWORK
Developers united in eZ Publish

Presenters

- **Donat** is owner of Webmanufaktur, a full service web agency in Switzerland. He works as projects manager, software architect and developer and likes thinking outside of the box. In the last year he has been involved in major eZ 5 projects.
- **Ekke** is a consultant with deep knowledge in eZ Publish 4 and 5, eZ Find / Apache Solr and with a faible for coming cutting edge web technologies. He is one of the organizers of the PHP Unconference since seven years.
- Members of CJW Network

Learnings from Real eZ Publish 5 Projects

- Last year we have presented our Cookbok...
- ... today you get a Ratatouille ;-)
- Let's share the experiences you and we made in the last year!

Agenda

Things we would like to discuss:

- Good Practice
- ez View Cache vs. HTTP Cache
- Debugging
- Pitfalls
- MultiSite Setup



Good Practice

Team up with a Symfony Crack

- To be honest: as eZ 4 developers, we are complete novices in eZ 5
- It's easier for a Symfony Crack to learn eZ than other way round
- Symfony community is hungry for a CMS, so watch out for new competition

Think in MVC

- A radical different thinking
- eZ 4 manged all together in the TPL -> the view implemented the logic (fetch)
- Symfony enforces a clean separation, from routing to the controller to the rendering of the view

Think in Bundles

- Split your application in different bundles (site specific, functional, ...)
- Reuse your code: create and maintain a Base Bundle with general functions

Creating bundles is easy, don't work in the DemoBundle ;-)

```
$ php ezpublish/console generate:bundle
```

BASH

Organize your config files

- The standard eZ installation is a mess...
- ... and the DemoBundle only slowly becoming a source of good practice

Our approach

- keep config in `ezpublish/config` as general as possible
- it should merely describe your environment, not the application
- move all site/function specific settings to the bundle

Keep ezpublish.yml small

ezpublish/config/ezpublish.yml

```
imports:
  - {resource: "@CjwExampleBundle/Resources/config/ezpublish.yml"}
```

YML

Keep these sections:

- doctrine (Database)
- ezpublish (Siteaccess Matching, Siteaccesses, Languages, Caching)
- stash

Can even be shorter - get inspiration from <https://github.com/lolautruche/metalfrance>

Note: prepending configuration does not work well with parameters

Config Files in Bundle

We keep them in a separate directory and name them as in good old eZ...

```
ExampleBundle
├─ Resources
│   └─ config
│       └─ ezpublish
│           ├── image.yml
│           ├── override.yml
│           ├── ezpublish.yml
│           ├── parameters.yml
│           ├── routing.yml
│           └── services.yml
```

Controllers

After several tries, we ended up with...

- Basically one controller per full view
- Separate controllers for navigation etc.
- Retrieve all needed data (location, children, ...)
- Prepare the data for easy processing in the templates
- Consider caching (TTL, X-Location Header)

Ways to Fetch Content

- `SearchService::findContent()`
- `SearchService::findLocations()`
- `LocationService::loadLocationChildren()`
- Legacy fetch functions

SearchService::findContent()

The only `SearchService` function you will find in `DemoBundle`

- returns full `content` objects with ALL attributes in ALL languages
- does not work well with multiple locations
- scales very badly
- no `asObjects` flag as in eZ 4
- fetching a content tree with 116 locations took 30 seconds
- most of the time is spent in manipulating the SQL result in PHP
- Another test: 24 hits, PHP array 44'880 rows with 39 elements each, highly redundant

<http://share.ez.no/blogs/donat-fritschy/searchservice-performance>

SearchService::findLocations()

Available from 2014.05 / 5.3 only

Roughly equivalent to good old `fetch('content', 'list')`

- returns `location` objects with `contentInfo` only
- fetching a content tree with 116 locations took < 1 second

LocationService::loadLocationChildren()

«Think of `LocationService::loadLocationChildren()` as primarily intended for administration interface.»

«If what it offers suits you for the frontend as well, great, but otherwise you will have to use `SearchService`. In other words, this method will not get filtering capabilities.»

(From a discussion with Petar)

<http://www.netgenlabs.com/Blog/Fetching-content-in-eZ-Publish-5-using-Search-service>

Legacy Fetch Functions

PHP

```
use eZFunctionHandler;
...
$mySearchResults = $this->getLegacyKernel()->runCallback(
    function () use ( $searchPhrase, $sort, $contentTypeIdentifiers )
    {
        // eZFunctionHandler::execute is the equivalent for a legacy template fetch function
        // The following is the same than fetch( 'ezfind', 'search', hash(...) )
        return eZFunctionHandler::execute(
            'ezfind',
            'search',
            array(
                'query'      => $searchPhrase,
                'sort_by'    => $sort,
                'class_id'   => $contentTypeIdentifiers
            )
        );
    }
);
```

TWIG Templates

- A generalized full and line view template for the easy stuff
- And again, basically one template per full view

How to handle children (sub items)?

- render directly in the template - the data is usually already there
- use `{{ render(controller('ez_content:viewLocation',
{ 'locationId': child.id, 'viewType': 'line' })) }}`

When to use ESI?

- nice concept, but quite a big overhead
- better suited for larger chunks

How to organize Templates?

The Symfony way...

```
views
├── Customer
│   └── CustomerDetail.html.twig
└── Product
```

The classic eZ way...

```
views
├── full
│   └── customer.html.twig
└── line
```

Two approaches, both valid. Follow your taste.



eZ View Cache vs. HTTP Caching

eZ View Caching

When the pagelayout is rendered, the `{ $module_result.content }` part will be replaced with the actual output. If view caching is enabled, the entire result of the module will be cached. This means that the contents of the "module_result" variable will be put into a cache file (...)

When a new version (...) of an object is published, the system will automatically clear the view cache for the following items:

- All published nodes of the object
- The parent nodes
- Related nodes (keywords, object relations)

<https://doc.ez.no/eZ-Publish/Technical-manual/4.x/Features/View-caching>

HTTP Expiration and Validation

The HTTP specification defines two caching models:

- With the **expiration model**, you simply specify how long a response should be considered "fresh" by including a Cache-Control and/or an Expires header. Caches that understand expiration will not make the same request until the cached version reaches its expiration time and becomes "stale";
- When pages are really dynamic (i.e. their representation changes often), the **validation model** is often necessary. With this model, the cache stores the response, but asks the server on each request whether or not the cached response is still valid. The application uses a unique response identifier (the Etag header) and/or a timestamp (the Last-Modified header) to check if the page has changed since being cached.

http://symfony.com/doc/current/book/http_cache.html

In Short (and much simplified...)

eZ View Cache caches **content** and **content fragements**

- Standard TTL is 2 hours
- Is purged on content modifications (with smart cache clearing rules)

Symfony's HTTP Cache caches **requests**

- eZ uses Expiration model by default
- Standard TTL is 60 seconds (86400 for tree menu!)
- Symfony Cache is purged from backend, **but only for ONE location**

Let's dive in a bit deeper...

ez 4 Cache Directives

Code from `ezpublish_legacy/kernel/private/classes/ezpkernelweb.php`

```
// send header information
foreach (
    eZHTTPHeader::headerOverrideArray( $this->uri ) +
    array(
        'Expires' => 'Mon, 26 Jul 1997 05:00:00 GMT',
        'Last-Modified' => gmdate( 'D, d M Y H:i:s' ) . ' GMT',
        'Cache-Control' => 'no-cache, must-revalidate',
        'Pragma' => 'no-cache',
        ...
    ) as $key => $value
)
{
    header( $key . ': ' . $value );
}
```

PHP

This guarantees that every request is handled by eZ

eZ 5 Cache Directives

Code from `vendor/ezsystems/demobundle/EzSystems/DemoBundle/Controller/DemoController.php`

```
// Setting HTTP cache for the response to be public and with a TTL of 1 day.  
$response = new Response;  
$response->setPublic();  
$response->setSharedMaxAge( 86400 );  
// Menu will expire when top location cache expires.  
$response->headers->set( 'X-Location-Id', $rootLocationId );  
// Menu might vary depending on user permissions, so make the cache vary on the user hash.  
$response->setVary( 'X-User-Hash' );
```

PHP

This effectively sets the Response free, out of the reach of eZ

Emulating eZ 4 Cache behaviour in eZ 5

This patch to `index.php` disables client and proxy caching without sacrificing the benefits of the Symfony HTTP cache. Use at own risk!

```
$response = $kernel->handle( $request );  
  
// Emulate eZ 4 cache control  
$response->headers->set( 'Cache-Control', 'no-cache, must-revalidate' );  
  
$response->send();  
$kernel->terminate( $request, $response );
```

PHP

Cache Recommendations

- Read the specifications
- Use Shared Caches with caution
- Set TTL values low enough
- Cave: `setTtl()` vs. `setClientTtl`
- For higher traffic sites, use Varnish
- <http://tools.ietf.org/html/rfc2616#page-74>
- https://www.mnot.net/cache_docs/



Debugging

Coping with blank screens

Blank screen, "503 Service not available"

- PHP errors (Syntax error, Memory, Outdated Autoloads, ...)
- Configuration errors (DB connection, ...)
- Switch to DEV mode for better debugging
- Check the log files

```
Apache/PHP Log  
ezpublish/logs/<env>.log  
ezpublish_legacy/var/log/*  
ezpublish_legacy/var/<siteaccess>/log/*
```

- Check write permissions on log files!

TwigBundle:Exception:error500.html.twig

- NEVER a Twig error!
- Caused by response 500 "Internal Server Error" and missing error template
- Checks as before

Twig Exception: Invalid variation "<variation>"

Caused by problems when accessing images

- Check if the file exists
- Check permissions on `ezpublish_legacy/var/<siteaccess>/storage`
- Check log files
- Clear cache

Class 'ezxFrmToken' not found

- Usually found with fresh installations involving legacy extensions
- Regenerate Autoloads

```
$ cd ezpublish_legacy  
$ php bin/php/ezpgenerateautoloads.php -e -p
```

BASH



Pitfalls

Avoid the traps...

Memory limit exceeded in DEV mode

- DEV mode takes a lot of memory
- Stash Logging is the worst
- Disable Stash Logging in ezpublish.yml

```
stash:  
  logging: false  
  caches:  
    default:  
      handlers:  
        - FileSystem  
      inMemory: true  
      registerDoctrineAdapter: false
```

YML

414 Request-URI Too Long

When doing subrequests, particularly ESI or Hinclude ones, current SiteAccess is transmitted in a serialized form, with its matcher. With a large number of configured SiteAccesses using Map\Host or Map\URI matcher (around 40, which is not uncommon for a multi-national, multi-lingual site) the URL can exceed the size of 8192 Bytes which most servers accept. As a result, the ESI/Hinclude call fails.

- Fixed in Version 5.3.3 (2014.07)
- <https://jira.ez.no/browse/EZP-23168>
- <https://github.com/ezsystems/ezpublish-kernel/pull/949>



Multi-Site/Multi-Repository Setup

Why a Multi-Site/Multi-Repository Setup?

- At CJW Network we have developed a multi-site/multi-repository setup for eZ Publish 4 several years ago
- This allows us to host many individual sites on a single eZ Publish installation

Advantages:

- Central site administration (site activation, cronjobs, ...)
- Easy deployment (update site extension with Subversion)
- Highly reduced maintenance costs (security patches, upgrades)
- Highly efficient use of hardware resources

Disadvantages:

- Some Kernel patches needed

Multi-Site/Multi-Repository Setup in eZ 5

First Approach (proven in production)

- Use different `ezpublish` directories to host the different sites

Second approach (under development)

- Use `CJW MultiSiteBundle`

TODO EKKE

Multi-Site-Setup (old) Directory structure

```
ezpublish                <-- not used
ezpublish_legacy
└ extension
  └ site_customer        <-- each customer has its own extension and database
  └ site_customertwo
└ var
  └ site_customer        <-- each customer has its own var directory
  └ site_customertwo
site_customer            <-- each customer has its own Symfony app
site_customertwo
src
└ CjwNetwork
  └ SiteCustomerBundle   <-- each customer has its own bundle
  └ SiteCustomertwoBundle
```

Multi-Site-Setup (old) Detail ezpublish_legacy

```
ezpublish_legacy
├── extension
│   ├── site_customer          <-- each customer has its own extension
│   ├── classes
│   ├── design
│   ├── modules
│   ├── settings
│   ├── site.ini
│   └── [DatabaseSettings]     <-- each customer has its own database
│       └── Database=database_site_customer
├── site_customertwo
│   ├── [...]
│   ├── settings
│   ├── site.ini
│   ├── [DatabaseSettings]
│   └── Database=database_site_customertwo
├── var
│   ├── site_customer          <-- each customer has its own var directory
│   └── site_customertwo
```

Multi-Site-Setup (old) Directory structure

```
site_customer
├─ autoload.php
├─ bootstrap.php.cache
├─ cache
├─ check.php
├─ config    <-- all yml files like ezpublish folder, (to improve)
│   └─ config.yml
│   └─ ezpublish.yml
├─ parameters.yml
├─ console
├─ logs
├─ phpunit.xml.dist
├─ Resources
├─ sessions
├─ SiteCjwbaseCache.php
├─ SiteCjwbaseKernel.php
└─ SymfonyRequirements.php
```

END TODO EKKE

Introducing CJW MultiSiteBundle

Although the first approach works fine, it has several drawbacks:

- Application code scattered at different places (site directory, bundle, legacy extension), hard to maintain in VCS, hard to deploy
- Redundancy in config files
- No global settings
- No central site activation/administration
- Goal: keep everything in one place!

CJW MultiSiteBundle Features

- Boots kernel and environment based on domain name mappings
- Handles local, staging and live domain names
- Allows for global activation of bundles
- Allows for global settings
- Provides a common console for all sites
- Caches domain name mappings
- Moves cache and log files away from the ezpublish folder
- more to come ...

cjwpublish Directory

The `cjwpublish` application directory sits next to the `ezpublish` directory.

```
cjwpublish
├─ config
│   ├── cjwpublish.yml          <-- defines active bundles
│   └── config.yml             <-- allows for global settings
├─ CjwPublishKernel.php        <-- inherits from CjwMultiSiteKernel.php
├─ CjwPublishCache.php         <-- inherits from CjwMultiSiteCache.php
└─ console
```

Symfony's app directory is back

Site Bundle Directory Layout

```
src
├── Cjw
│   └── SiteExampleBundle
│       ├── app
│       │   ├── config
│       │   │   ├── cjwpublish.yml      <-- contains domain mappings
│       │   │   ├── config.yml
│       │   │   ├── ezpublish.yml
│       │   │   └── ...
│       │   ├── CjwSiteExampleKernel.php <-- inherits from CjwPublishKernel.php
│       │   ├── CjwSiteExampleCache.php <-- inherits from CjwPublishCache.php
│       │   └── Controller
│       │       └── ...
```


Caveats

Adjustments needed in `config.yml` to reflect different relative location of kernel

```
assetic:
  ...
  read_from:      %kernel.root_dir%/../..../web
  write_to:       %kernel.root_dir%/../..../web
  ...
ez_publish_legacy:
  ...
  root_dir: %kernel.root_dir%/../..../ezpublish_legacy

parameters:
  ezpublish.kernel.root_dir: %kernel.root_dir%/../..../vendor/eZsystems/ezpublish-kernel
```

YML

More problems of this kind expected!

Project Status

- Currently in private Beta, not yet released
- Ideas and Feedback welcome
- Public Beta in October
- info@cjw-network.com
- <https://github.com/cjw-network/MultiSiteBundle>

Ressources

Slides as PDF

- see `src/Cjw/SummerCampBundle/Resources/doc`
- <https://github.com/cjw-network/SummerCampBundle/Resources/doc>

Slides (Source)

- <https://github.com/dfritschy/cjw-summercamp-slides>

CJW MultiSiteBundle

- <https://github.com/cjw-network/MultiSiteBundle>
- info@cjw-network.com



<http://vote.netgenlabs.com/>

Please Vote!

<Thank You!>

Ekkehard Dörre

<http://share.ez.no/community/profile/7431>

[@ekkeD](#)

<http://www.coolscreen.de>



Donat Fritschy

<https://github.com/dfritschy>

<http://share.ez.no/community/profile/10451>

[@webmanufaktur](#)

<http://www.webmanufaktur.ch>